

Discussion: Four Questions

- 1. Are language models intelligent?**
- 2. Can humans be trusted with this tool?**
- 3. Is AI a net good for society?**
- 4. Are we doomed?**

My Answers

1. Are language models intelligent?

- Yes. These models “understand” things.
- How do we understand things?

2. Can humans be trusted with this tool?

- Clearly not. But stopping the train is unrealistic.

3. Is AI a net good for society?

- All tools are neutral. It's how we use them that matters.
- Nothing new about humans using tools for bad purposes.
- We've been on this path for a while (internet, smart phones, social media).
- The solution to most problems is education.

4. Are we doomed?

- Possibly, but given my answer to 1, I'm not sure that's all bad. :)
- Geoff Hinton's view